

iNLP Assignment-3: Static Word Embeddings [2023201044]

Code Explanation:

1. Brown Corpus Preprocessing (brown.py)

- `preprocess_brown()`:
 - Loads the Brown Corpus from the nltk library.
 - Tokenizes sentences and words.
 - Converts words to lowercase and removes punctuation.
 - Returns a processed corpus where each sentence is represented as a list of words.

2. Continuous Bag of Words (CBOW) Model (cbow.py)

- Corpus Processing:
 - Calls `preprocess_brown()` to load the Brown corpus.
 - Loads a precomputed vocabulary (`sorted_vocab.json`).
- Vocabulary Building (`build_vocab`):
 - Assigns indices to words for encoding.
 - Uses frequency-based sorting.
- Training Data Generation (`generate_training_data`):
 - Creates context-target pairs for CBOW training.
- CBOWNegativeSampling Model:
 - Defines an embedding layer for words and a separate one for outputs.
 - Implements forward propagation using negative sampling.
- Negative Sampling (`get_negative_samples`):
 - Selects negative words using a weighted probability distribution.
- Training and Hyperparameter Tuning:
 - Optimizes hyperparameters (context size, negative samples, embedding dimensions).
 - Uses early stopping based on WordSim-353 evaluation.
 - The **hyperparameters** used are:
 - half-context window size (2, 3, 5)

- number of negative samples (2, 5, 7)
- embedding size (300)
- Adam optimizer with 0.001 learning rate
- Evaluation (evaluate_model):
 - Computes word similarity and evaluates embeddings.

3. Skip-gram Model (skipgram.py)

- Vocabulary Construction:
 - Builds vocabulary from the Brown Corpus.
 - Uses frequency-based sorting and filtering.
- Training Data Preparation (generate_training_data):
 - Generates (target, context) word pairs.
- Dataset Class (Word2VecDataset):
 - Implements a PyTorch dataset for Skip-gram training.
 - Generates negative samples dynamically.
- Skip-gram Model (SkipGramBinaryClassifier):
 - Uses embeddings for target and context words.
 - Computes similarity scores with a dot product.
- Training Routine (train_word2vec):
 - Loads dataset and optimizes word embeddings.
 - Uses early stopping based on WordSim-353 evaluation.
 - The **hyperparameters** used are:
 - half-context window size (2, 3)
 - number of negative samples (3, 5, 7)
 - embedding size (300)
 - Adam optimizer with 0.001 learning rate

4. Singular Value Decomposition (SVD) Model (svd.py)

- Corpus Processing and Vocabulary Creation:
 - Filters words with minimum frequency of 3 (replacement with <UNK>).
 - Creates a co-occurrence matrix.
- Co-occurrence Matrix (build_co_occurrence_matrix):

- Constructs a word-word matrix using a context window.
- Optionally applies distance-based weighting and stopwords removal.
- SVD Decomposition:
 - Applies Singular Value Decomposition (SVD) to extract word vectors.
 - Normalizes embeddings for better performance.
- Hyperparameter Optimization:
 - Tests different window sizes and embedding dimensions.
 - Evaluates model using WordSim-353.
 - The **hyperparameters** used are:
 - half-context window size (2, 4, 6)
 - number of negative samples (2, 5, 7)
 - embedding size (100, 200, 300)
 - apply weighting (closer words carry more weighting in the co-occurrence matrix) (False, True)
 - remove stopwords (False, True)

5. Evaluation Scripts (wordsim.py & wordsim_score.py)

- Loading and Evaluating Embeddings (evaluate_model):
 - Loads embeddings from trained models.
 - Computes cosine similarity for word pairs.
 - Evaluates using Spearman's Rank Correlation against human scores.
- Command-line Execution (wordsim.py):
 - Loads pre-trained embeddings from user input.
 - Calls evaluate_model() to compute similarity scores.

Instructions for file execution:

1. To train any model, run the corresponding script as:
`python model_name.py`
2. To compute cosine similarities and Spearman's correlation coefficient, run the script as:
`python wordsim.py <embedding_path>`

Results:

The pre-trained models can be found [here](#).

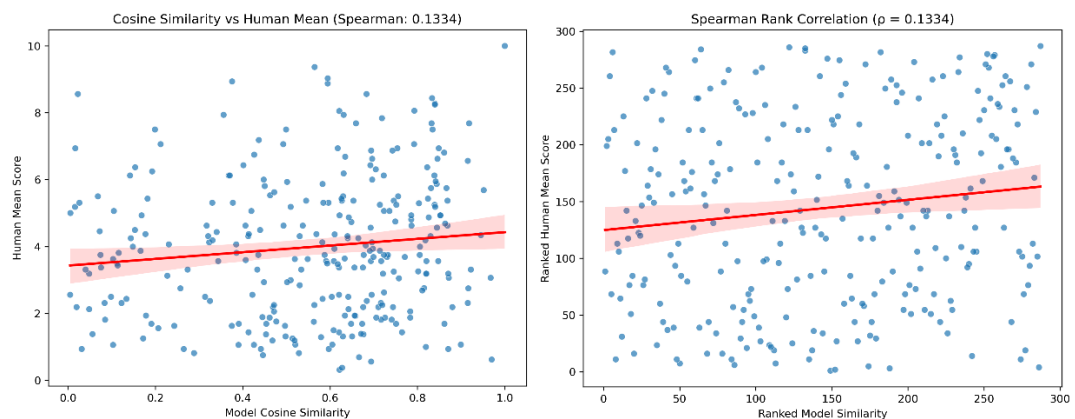
1. SVD:

a. Best hyper-params:

weighting=False, remove_stopwords=True, window_size=4, embedding_dim=100

The following are the first 5 and last 5 rows of results from hyper-parameter tuning.

Weighting	Remove Stopwords	Window Size	Embedding Dim	Spearman Correlation
False	True	4	100	0.133439
True	True	4	100	0.127971
False	True	4	300	0.126675
True	True	6	100	0.120008
False	True	6	100	0.118182
False	False	4	200	0.003220
False	False	4	100	-0.007677
False	False	6	300	-0.014319
False	False	6	200	-0.037897
False	False	6	100	-0.052195



b. Wordsim results:

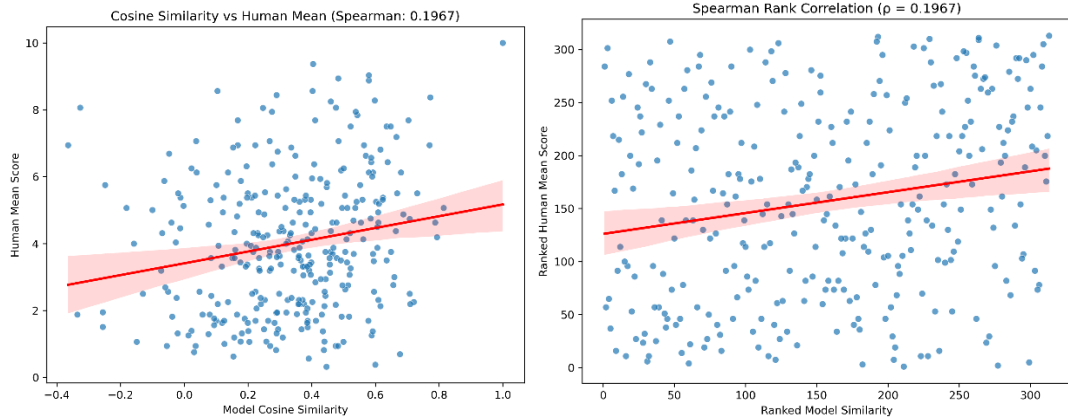
Word 1	Word 2	Human (Mean)	Model Similarity
admission	ticket	5.536	0.4626487195491791
alcohol	chemistry	4.125	0.3229874074459076
aluminum	metal	6.625	0.8454661965370178
announcement	effort	2.0625	0.471233069896698
announcement	news	7.1875	0.4363973140716553

c. Spearman's rank correlation = 0.1334

2. CBOW:

a. Best hyper-parameters (20 epochs):

CONTEXT_SIZE=2, NEGATIVE_SAMPLES=7, EMBEDDING_DIM=300



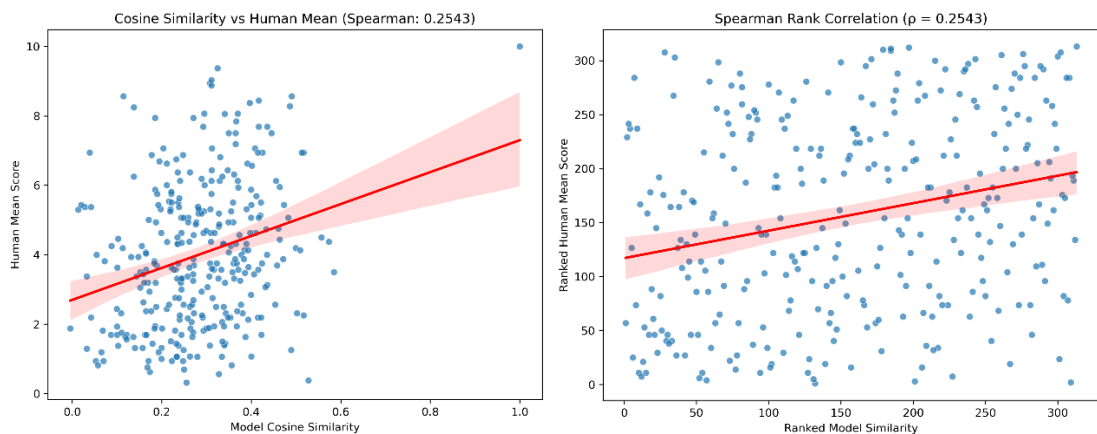
- b. Note that the weights in the first layer are used as embeddings. Alternately, the weights in the second layer or the averaged weights could be used as embeddings.
- c. Wordsim results:

Word 1	Word 2	Human (Mean)	Model Similarity
admission	ticket	5.536	0.2536108195781708
alcohol	chemistry	4.125	0.5613850355148315
aluminum	metal	6.625	0.569373369216919
announcement	effort	2.0625	0.22620703279972076
announcement	news	7.1875	0.6652630567550659

- d. Spearman's rank correlation = **0.1967**

3. SGNS:

- a. Best hyper-params (10 epochs):
CONTEXT_SIZE=2, NEGATIVE_SAMPLES=7, EMBEDDING_DIM=300



- b. Wordsim Results:

Word 1	Word 2	Human (Mean)	Model Similarity
admission	ticket	5.536	0.21598714590072632
alcohol	chemistry	4.125	0.38459086418151855
aluminum	metal	6.625	0.39151731133461
announcement	effort	2.0625	0.2668076753616333
announcement	news	7.1875	0.36513423919677734

- c. Spearman's rank correlation = **0.2543**

Analysis:

The Skip-gram model with Negative Sampling (SGNS) outperforms both CBOW and SVD in the WordSim-353 task. While SVD provides a useful approximation, it struggles with less frequent words due to its reliance on a co-occurrence matrix. CBOW, though faster, is less effective in preserving rare word relationships compared to SGNS..

1. SVD:

- Benefits: Captures global word co-occurrence relationships; computationally efficient for **offline processing**.
- Limitations: Less effective for rare words; requires **significant memory** for storing the full co-occurrence matrix.

2. CBOW:

- Benefits: **Fast training** due to averaging of context vectors and fewer updates; learns representation of more frequent words effectively.
- Limitations: Loses positional context information, struggles with rare words compared to SGNS.

3. SGNS:

- Benefits: **Best performance** on WordSim-353, better representation of **less frequent** words; produces higher quality embeddings, capturing subtle semantic nuances.
- Limitations: Requires **more training time** and computational power due to dynamic negative sampling.